

Tutorial on SimpleDB

Jia Wang
Oct 11, 2016

client

Client program

server

Server startup
program

remote

server

Remote driver

SimpleDB
inst.

Simple driver

Buffer manager

buffer

Metadata
manager

metadata

File manager

file

Log manager

log

Transaction

tx

parse

Parser

Planner

*planner
opt*

Scan/Plan

query

Record file

record

Block/Page

file

Planner

- Input: SQL string
- Output: a QueryPlan
- Calls parser to extract **query data**
--tables, fields, predicates... from
input
- Plans the **order** of relational
operations and **which** algorithms to
be used-- constructing a plan tree

Running example

- Two tables
 - Students table: (SID, SName, MajorId, Age)
 - Dept table: (DID, DName)
- Query
 - ***SELECT*** SName, DName
FROM Students, Departments
WHERE MajorId = DID ***AND*** Age > 20;

SID	SName	MajorId	Age
1	Alice	10	18
2	Bob	20	20
3	Carl	30	22

DID	DName
10	ECE
20	Math
30	CS

Starting server

- Start-up program creates **SimpleDB instance**
 - See: [simplifiedb/server/Startup.java](#)
- Creates a (remote) **driver** to drive the SimpleDB inst.
 - See: [simplifiedb/remote](#)
- The SimpleDB instance includes various **managers**
 - Metadata manager
 - E.g., how many records each table contains
 - File/buffer manager
 - The files/buffers stored by this database
 - ...
 - See: [*Mgr.java](#) in [simplifiedb/server](#), [simplifiedb/metadata](#), [simplifiedb/file](#), [simplifiedb/buffer](#)

Client program

- Creates a **driver** that connects to the server
- Creates a **Statement** object from the driver
- The statement object will **execute** user query in string format

```
Driver d = new SimpleDriver();
conn = d.connect("jdbc:simpledb://localhost",
null);

Statement stmt = conn.createStatement();

String qry = "select dname, age"
            + "from student, dept "
            + "where DID = MajorId AND age >= 20 ";
ResultSet rs = stmt.executeQuery(qry);

while (rs.next()) {
    String dname = rs.getString("dname");
    int age = rs.getInt("age");
    System.out.println(dname + "\t" + age);
}

rs.close();
```

Parsing

- When the `stmt.executeQuery(q)` is called, a **Planner** will be created to *actually* execute the query, and the result is returned as a `ResultSet` object.
 - Meanwhile, a **Transaction** object, associated with the query, will also be created.
 - See: [simplifiedb/remote](#), [simplifiedb/tx](#)
- The planner calls **Parser** to parse the query string into a `QueryData` instance, which contains the **structured data** of a query
 - The list of *tables* being queried
 - The *predicate* for selection
 - The list of *fields* for projection
 - See: [simplifiedb/parser](#)

Query data parsed from query:

SELECT SName, DName

FROM Students, Dept

WHERE MajorId = DID **AND** Age >= 20;

Fields

- SName
- DName

Tables

- Students
- Dept

Predicate

- (MajorId = DID) conj. (Age >= 20)

Terms

- MajorId = DID
- Age > 20

Expressions

- MajorId
- DID
- Age
- 20

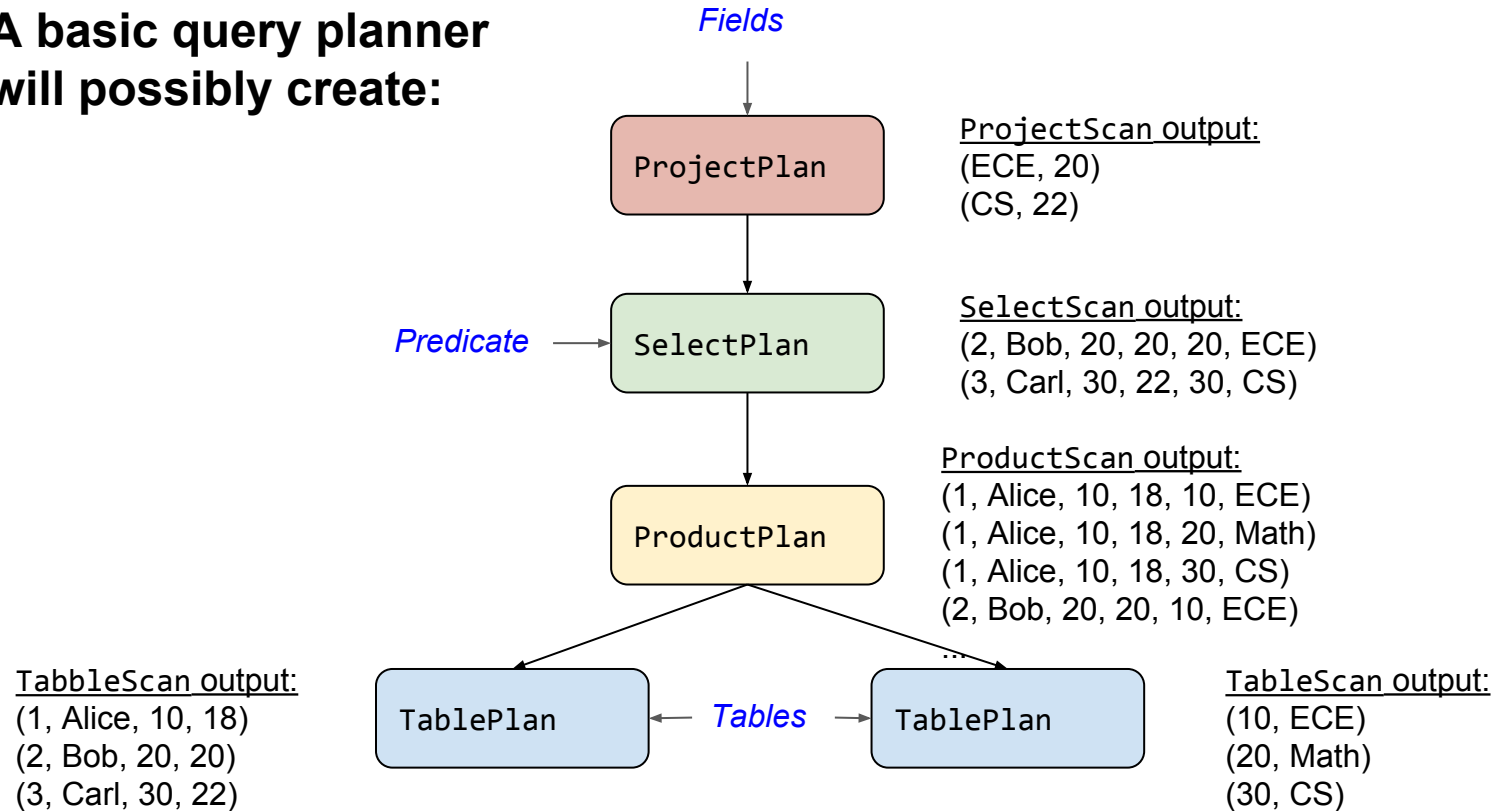
Query processing and optimization

- Planner creates a **QueryPlan** from **QueryData**, which is a **tree** of **Plan**'s.
 - Each relational operation can be processed in different ways, each way is represented as a **Plan**.
 - E.g, **IndexJoinPlan**, **HashJoinPlan**, **SelectPlan**, **IndexSelectPlan**.
 - Each **non-leaf Plan** contains **sub Plan**'s, while a **leaf Plan** must be a **TableScan**.
 - E.g., The Plan for a *join* operation contains two sub Plan's, one for scanning the left-hand table and one for the right-hand table.
 - The Planner can use **optimization** techniques to create an cost-efficient query plan tree.
 - Folders: [simplifiedb/query/*Plan.java](#), [simplifiedb/planner](#), [simplifiedb/opt](#)

Query processing and optimization

- Each `Plan` object is associated with a `Scan` object, which actually computes/returns the output of the operation **record-by-record**.
 - `getInt()`, `getString()`, `hasField(fld)`, etc.
 - If parent `Plan` `p` wants to *read output* from its sub-plan `p'`, then `p` needs to call `p'.open()` to get a `Scan` `s`, and get values of a integer field by `s.getInt()`, and moves to *next* output by `s.next()`
 - See: [simplifiedb/query/*Scan.java](#)

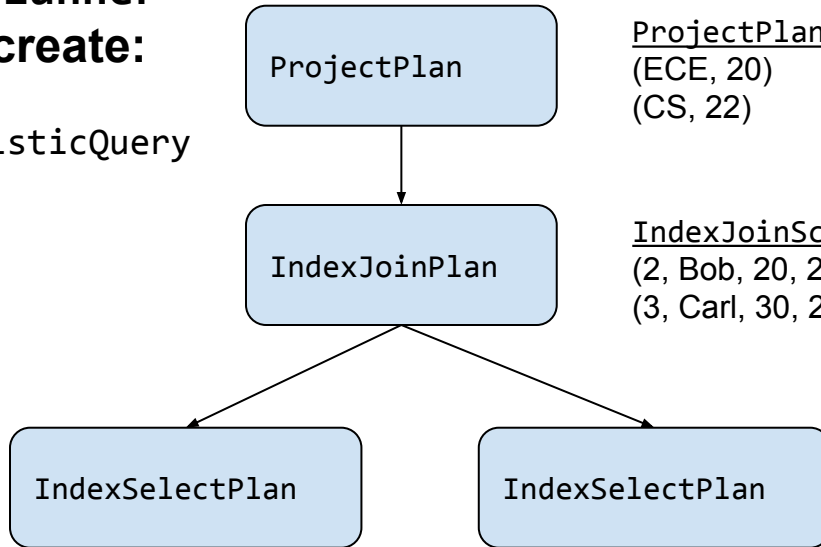
A basic query planner will possibly create:



A better query Planner might possibly create:

(e.g.,
simplifiedb/opt/HeuristicQuery
Planner.java)

Assuming table
“students” has an
index on “age”...



ProjectPlanScan output:
(ECE, 20)
(CS, 22)

IndexJoinScan output:
(2, Bob, 20, 20, ECE)
(3, Carl, 30, 22, CS)

IndexSelectScan output:
(2, Bob, 20, 20)
(3, Carl, 30, 22)

IndexSelectScan output:
(10, ECE)
(20, Math)
(30, CS)

Assuming table “dept”
has an index on
“DID”...

How a basic query planner creates a query plan:

```
public Plan createPlan(QueryData data, Transaction tx) {  
    //Step 1: Create a plan for each mentioned table or view  
    List<Plan> plans = new ArrayList<Plan>();  
    for (String tblname : data.tables()) {  
        String viewdef = SimpleDB.mdMgr().getViewDef(tblname, tx);  
        if (viewdef != null)  
            plans.add(SimpleDB.planner().createQueryPlan(viewdef, tx));  
        else  
            plans.add(new TablePlan(tblname, tx));  
    }  
  
    //Step 2: Create the product of all table plans  
    Plan p = plans.remove(0);  
    for (Plan nextplan : plans)  
        p = new ProductPlan(p, nextplan);  
  
    //Step 3: Add a selection plan for the predicate  
    p = new SelectPlan(p, data.pred());  
  
    //Step 4: Project on the field names  
    p = new ProjectPlan(p, data.fields());  
    return p;  
}
```

Record scanning

- Each **leaf** in the query plan tree will be a **TableScan** plan, which actually reads **records** from the storage.
 - See [simplifiedb/query/TableScan.java](#), [simplifiedb/record](#)
- The scanning can be either a simple **linear** scan or **index** scan.
 - An index scan will contain a **TableScan** object which returns records specified by the index
 - See [simplifiedb/index](#), [simplifiedb/query](#)
- The records will be stored in disk **blocks**.
 - Each disk block is referred to by a **Block** instance, and the actual content of a block is stored by a **Page** instance. So **Block** and **Page** instances are always used together.
 - See [simplifiedb/file](#)

Using indexes

- Create an index *just like* creating a table
 - Specifying the table and the fields.
 - Managed by an IndexManager (similar to TableManager).
- When inserting records into the table, **also** *insert into the index*.
 - See [simplifiedb/index/planner/IndexUpdatePlanner.java](#)
- When using index for table scanning
 - Both a TableScan and the Index will be used
 - The Index will return the **pointer** to a record, while the TableScan will move to that pointer to return the record **content**.
 - See [simplifiedb/index/query/*.java](#)

Example project

- To add a **new join algorithm**, which parts of the code will be touched?
 - **Plan**: create a plan class for the new join operation using your algorithm
 - It *estimates the **cost*** of join using the new algorithm
 - It *creates the **Scan object*** for your new algorithm.
 - It *determines the **schema*** of the output of this operation.
 - See: [simplifiedb/query/*Plan.java](#)
 - **Scan**: a new scan class corresponding to the new plan, which *actually* computes/returns the output of the join using your algorithm *one-by-one*.
 - The new plan will know how to create this new scan
 - See: [simplifiedb/query/*Scan.java](#)

Example project

- **Planner:** now the planner *knows* that your new join operation exists, and would consider *using* it when proper when it constructs the plan tree (e.g., based on cost estimation).
 - Make your own Planner or modify existing planners.
 - See: [simplifiedb/opt/*Planner.java](#), [simplifiedb/planner/*Planner.java](#).
 - Change [simplifiedb/server/SimpleDB.java](#) to switch to another planner.
- **Parser:** modify Parser if the SQL needs be extended.
 - See: [simplifiedb/parser](#).